

Franklin's leader election algorithm for bidirectional rings

Introduction

Franklin's algorithm discussed here performs leader election in bidirectional rings with asynchronous communication. Each process (node) has a unique identifier id known to itself. The pessimistic number of messages is at most $O(n \log_2 n)$.

The algorithm - description

At any time of the algorithm's execution, each process is either active, passive or a leader. There are either no leaders or exactly one leader at any given moment. Before an overview, let's define 'active neighbours'. By 'left/right active neighbour of node a ' we mean the closest active node when traversing the ring in the direction specified by a 's left/right link. Nodes may have different understanding of which side is 'left' and which 'right'; this is not an issue.

The algorithm is split into rounds. In each round an active process a sends two messages with its id to left and right active neighbours and receives two messages with $idLeft$ and $idRight$ from these nodes. If either of received identifiers is greater than a 's id , then it becomes passive. If $idLeft == id$ or $idRight == id$, a becomes the leader. If neither of above conditions is met, a remains active.

Passive node a simply forwards messages in from either direction in the same direction. If a message is *final*, then a sets message's value as leader's identifier and terminates. It will become clear that an active node cannot receive a *final* message. A leader node only sends a single *final* message and terminates when it comes back (forwarded by other nodes, all passive).

The algorithm - pseudocode

Process a maintains:

- id_a — a unique identifier
- $leaderId$ — leader's id , must be set before termination
- $state_a \in \{active, passive, leader\}$

Message format:

- id — original sender's id
- $isFinal$ — true for the single message sent by leader and then forwarded; false otherwise

Procedures

Initialization: all processes are active and know their id 's

Passive process: Upon receipt of a message msg , passes on the message. If $msg.isFinal$ is true, sets $leaderId$ to $msg.id$ and terminates.

Active process: In each phase a :

- sends message with id_a to its left and right neighbours,
- receives messages with id_{left} and id_{right} from its active neighbours;
- if $id_{left} > id_a$ or $id_{right} > id_a$, then a becomes passive,
- if $id_{left} == id_a$ or $id_{right} == id_a$, then a becomes a leader

Leader process: Leader process a sets $leaderId = id_a$ and sends message with id_a and $isFinal$ set to *true* to its left or right neighbour. When it receives the message back, terminates.

Correctness

Correctness of Franklin's algorithm stems from the following observations:

- If there is a single *leader* node and other nodes are passive, all nodes will terminate with correct *leaderId* set. This is clear from the protocol.
- If there is a single *leader* node, all other nodes are passive. This is because in order to become a leader, node *a* has to receive a message with its *id*. Such message can only be sent by *a* (ad identifiers are unique). It can be received back only when it traversed entire ring, and this implies that all nodes other than *a* are passive.
- There can't be two leader nodes. This fact follows from the above point.
- If there are at least two active processes in a given round, at least one of them must become passive at the end of this round. If this was not true, each of the active nodes would have *id* bigger than its neighbour identifiers, which is impossible.
- The algorithm is deadlock-free. This is clear from the protocol.

Complexity analysis

Let n be the size of the ring. In each round, at least half of the active nodes become passive, as in order to remain active a node must have *id* greater than its active neighbours identifiers. Therefore, after at most $FLOOR(\log_2 n)$ steps there is only one active process. When only one process *a* is active, it will send two messages which will be forwarded by the other nodes and come back to *a* - $2n$ message passes. Then *a* becomes leader and after the next n passes of the final message, the algorithm terminates. As there were exactly $2n$ message passes in any round with $k \geq 2$ active nodes, we conclude that the time complexity is:

$$2n * FLOOR(\log_2 n) + 3n \in O(n \log_2 n)$$

It is stated in [1] that average number of messages passed was $O(n)$ in sample tests.

References

- [1] Randolph Franklin. "On an improved algorithm for decentralized extrema finding in circular configurations of processors". In: *Communications of the ACM* 25.5 (1982), pp. 336–337.