

# Ben-Or asynchronous Byzantine consensus algorithm

## Introduction

The idea of the consensus problem is to have a system of  $N$  processes,  $t$  of which are faulty. Each correct process starts with either 0 or 1, and after running the algorithm, all correct processes must end up with either 0 or 1 based on the following rules:

- All correct processes end up with the same value.
- If all correct processes start with the same value  $v$ , then all correct processes must end with  $v$ .
- All correct processes must terminate (with probability 1).

Ben-Or algorithm allows the network to be asynchronous and does not restrict the malevolence of faulty processes.

## Protocol

Let  $N$  - number of all processes,  $t$  - number of faulty processes. It is assumed that  $t < \frac{N}{5}$ .  
Let's bring up the pseudo code[1]:

**Process P:** Initial value  $x_p$ .

**step 1:** Set  $r := 1$ .

**step 0:** Send the message  $(1, r, x_p)$  to all the processes.

**step 2:** Wait till messages of type  $(1, r, *)$  are received from  $N - t$  processes. If more than  $(N - t)/2$  messages have the same value  $v$ , then send the message  $(2, r, v, D)$  to all processes. Else send the message  $(2, r, ?)$  to all processes.

**step 3:** Wait till messages of type  $(2, r, *)$  arrive from  $N - t$  processes.

(a) If there are at least  $t + 1$  D-messages  $(2, r, v, D)$ , then set  $x_p := v$ .

(b) If there are more than  $(N + t)/2$  D-messages then **decide**  $v$ .

(c) Else set  $x_p = 1$  or 0 each with probability  $\frac{1}{2}$ .

**step 4:** Set  $r := r + 1$  and go to step 1.

Processes locally keep track of current round number, and store current binary value.

Within each round they go through 2 stages, messages from each round from now on will be called type 1 or 2 messages.

## Correctness

Let's start with proving the following lemma:

**Lemma 1.** *If all correct processes start with the value  $v$ , then within one round they will all decide  $v$ .*

*Proof.* Each process broadcasts message  $(1, 1, v)$ .

Process then waits for  $N - t$  messages.

At most  $t$  messages came from faulty processes, which means that at least  $N - 2 \cdot t$  messages are correct. Because  $N - 2 \cdot t > \frac{N+t}{2}$ , all the correct processes send message  $(2, 1, v, D)$ .

Among  $N - t$  accepted type 2 messages at most  $t$  are incorrect, which means that step 3(b) will not execute.

Again, because at least  $N - 2 \cdot t$  of type 2 messages are correct, then more than  $\frac{N+t}{2}$  of them will have the same value ( $v$ ). That means that every process will **decide**  $v$  this round.  $\square$

**Lemma 2.** *If process sets  $v$  in step 3(a) in round  $r$ , then it can't set  $\neg v$  in the same round.*

*Proof.* Let's say that process saw  $\geq t+1$  messages  $(2, r, 0, D)$  and  $\geq t+1$  messages  $(2, r, 1, D)$ . Let  $A_x$  be equal to the number of messages  $(2, r, x, D)$  originated from correct processes with  $x = 0, 1$ . Of course  $A_0 + A_1 \geq t + 2$ .

Every process responsible for  $A_x$  saw more than  $\frac{N-t}{2}$  messages  $(1, r, x)$  from correct process. But that means that there are more than  $N - t$  correct processes, which is a contradiction.  $\square$

Then let's look at the next lemma, which states that the processes will decide with at most 2 round window.

**Lemma 3.** *If for some round  $r$ , some correct process decides  $v$  in step 3(b), then all other correct processes will decide  $v$  within the next round.*

*Proof.* Let  $P$  be the process that has decided on  $v$ . In order to **decide**  $v$ ,  $P$  must have received more than  $\frac{N+t}{2}$  messages  $(2, r, -, D)$ . That means that more than  $\frac{N-t}{2}$  correct processes sent that message. Let  $A_x$  be the number of  $(2, r, x, D)$  messages from correct processes for  $x = 0, 1$ .

Because  $P$  has decided  $v$ , it follows that  $A_{\neg v} \leq t$ . On top of that  $A_v + A_{\neg v} > \frac{N-t}{2}$ .

Simple calculations show that  $A_v \geq t+1$ . It means that every other process will set  $v$  in step 3(a). From *Lemma 2* we know that it will be the only value that they will set.

It follows that in the next round every correct process will start with  $v$ , so *Lemma 1* can be applied.  $\square$

## Time complexity

In each round processes have probability of at least  $2^{-(N-t)+1}$  of all setting the same value, which using *Lemma 1* means that the run algorithm would end next round. That means that expected number of rounds is bounded by  $O(2^{N-t})$ .

The following theorem shows tighter bound under stricter assumptions:

**Theorem 1.** *[1] If  $t = O(\sqrt{N})$  then the expected number of rounds to reach agreement in this protocol is constant, i.e., it does not depend on  $N$ .*

## References

- [1] Michael Ben-Or. Another Advantage of Free Choice: Completely Asynchronous Agreement Protocols. In *Proceedings of the second annual ACM symposium on Principles of distributed computing*, pages 27-30. ACM, 1983, doi: 10.1145/800221.806707.